

# SENTRY

 [admin](#)

 [July 23, 2019](#)

 [Leave a comment](#)

 [Edit](#)

Sentry is an open-source **error tracking tool** that helps you monitor and fix crashes in real time.

Sentry captures data by using an SDK within your application's runtime. These are platform specific and allow Sentry to have a deep understanding of both how your application works.

Design → Code & Review → Testing/QA → CI/CD → Issue discovery → Investigation → Remediation

Plays a crucial role in [post-deployment](#) process of discovery, investigation, and remediation.

Sentry service level configuration for [spring](#). It is configured as part of NextService jar.

### MICROSERVICES:

Sentry for microservices is configured as part of **NextServiceArchetype**. It involves creation of

1. **SentryServletInitializer** bean
2. **SentryExceptionHandler** bean

They are both initialized as part of **SentryConfiguration** component in archetype.

To create a new dashboard for a project :

1. From the admin panel, create a new dashboard
2. Get the client dsn for the dashboard
3. Assign this value to “**Sentry.project.url**” in corresponding properties file of the service.

## SETUP:

Minimum services to be running before, setting up sentry :

1. PostgreSQL
2. Redis

Sentry also supports horizontal scaling, i.e running on a cluster of machines.

Things to consider before setting up of sentry :

1. TTL for events (Storage of historical data)
2. Average event throughput
3. How many events get grouped together

Always restart Sentry services **web, worker, cron** after [upgrading](#).

Two essential config files required for sentry configuration :

1. Config.yml [Configure core sentry attributes]
2. Sentry.conf.py [Python file is loaded after all configuration files are loaded]

[Configurable parameters](#) as part of **config.yml**

## GENERAL

1. system.admin-email [reported to sentry team as part of beacon]
2. system.url-prefix [Reference in UI as well as outbound notifications]

3. `system.secret-key` [Used for session signing]

## REDIS

1. `redis.clusters` [referenced by cache, digests, and TSDB backends]

## MAIL

---

[Configurable parameters](#) as part of `sentry.python.conf`

## WEB SERVER

1. `SENTRY_WEB_HOST`
2. `SENTRY_WEB_PORT`,
3. `SENTRY_WEB_OPTIONS` (dictionary of additional configuration options to pass to uwsgi.)]

## SMTP SERVER

## BEACON

1. `SENTRY_BEACON` – [allows communication with remote sentry servers]
- 

Sentry provides an abstraction called ‘filestore’ which is used for storing files. The default backend stores files on the local disk in `/tmp/sentry-files` which is not suitable for production use. In **`config.yml`**, following options can be placed for attaching a file storage system.

`filestore.backend: ‘S3’`

filestore.options:

access\_key: '...'

secret\_key: '...'

bucket\_name: '...'

---

## ASYNCHRONOUS WORKERS :

1. Sentry comes with a built-in queue to process tasks in a more asynchronous fashion.

For example when an event comes in instead of writing it to the database immediately, it sends a job to the queue so that the request can be returned right away, and the background workers handle actually saving that data.

2. Sentry also needs a cron process.

Above operations can be managed by using [supervisor](#).

Sentry currently supports two primary backends : Redis, RabbitMq

Drawback with Redis : all pending work must fit in memory

Eg: BROKER\_URL = "redis://:password@localhost:6379/0"

## WRITE BUFFERS :

Sentry manages database row contention by buffering writes and flushing bulk changes to the database over a period of time. This is extremely helpful if you have high concurrency, especially if they're frequently the same event.

For example, if you happen to receive 100,000 events/second, and 10% of those are reporting a connection issue to the database (where they'd get grouped together),

enabling a buffer backend will change things so that each count update is actually put into a queue, and all updates are performed at the rate of how fast the queue can keep up.

Sentry provides a service to store **time-series data**. Primarily this is used to display aggregate information for events and projects, as well as calculating (in real-time) the rates of events.

It will use the default [Redis cluster](#) configured unless specifically configured to use another one.

Historically, SENTRY\_CONF or --config was pointed directly to your sentry.conf.py. Now, SENTRY\_CONF should be pointed to the parent directory that contains both the python file and the yaml file. sentry init will generate the right structure needed for the future

## THROTTLES AND RATE LIMITING:

Too much inbound traffic without a good way to drop excess messages.

Event Quota : throttling workloads in Sentry involves setting up event quotas.

It provides rate limiting with a wide range of systems. [Refer](#)

Sentry relies heavily on [queues](#), monitoring them is essential.

Sentry provides several ways to monitor the [system status](#).

Sentry has a [CLI](#) and provides different options to achieve those which are not possible with UI.

## CURRENT SETUP:

## GENERAL ISSUES:

1. Refer [FAQ](#) first

We are running SENTRY as a docker container in an EC2 machine.

Inside the EC2 machine we have PostGreSQL and Redis configured necessary for the functioning of Sentry.

Following are the docker commands to run SENTRY :

```
docker run -it --rm -e SENTRY_URL_PREFIX='http://54.152.37.40:8090' -e
SENTRY_SECRET_KEY='+i78z!+x*t*(e#qrpq#yojt5fr#td0ztv#sd!&uiwg!gcuy706'
-e SENTRY_REDIS_HOST='10.0.1.46' -e
SENTRY_POSTGRES_HOST='10.0.1.46' -e SENTRY_POSTGRES_PORT='5432' -
e SENTRY_DB_NAME='nextsentry' -e SENTRY_DB_USER='sentryapp' -e
SENTRY_DB_PASSWORD='Gh75Dc' -e
SENTRY_EMAIL_HOST='mx.learnnext.com' -e SENTRY_EMAIL_PORT='25' -e
SENTRY_SERVER_EMAIL='system@nexteducation.in' sentry upgrade
```

#————— MY SENTRY —————

```
docker run -d --restart always --name my-sentry -e
SENTRY_URL_PREFIX='http://54.152.37.40:8090' -e
SENTRY_SECRET_KEY='+i78z!+x*t*(e#qrpq#yojt5fr#td0ztv#sd!&uiwg!gcuy706'
-e SENTRY_REDIS_HOST='10.0.1.46' -p 8090:9000 -e
SENTRY_POSTGRES_HOST='10.0.1.46' -e SENTRY_POSTGRES_PORT='5432' -
e SENTRY_DB_NAME='nextsentry' -e SENTRY_DB_USER='sentryapp' -e
SENTRY_DB_PASSWORD='Gh75Dc' -e
SENTRY_EMAIL_HOST='mx.learnnext.com' -e SENTRY_EMAIL_PORT='25' -e
SENTRY_SERVER_EMAIL='system@nexteducation.in' sentry
```

#————— SENTRY CRON —————

```
docker run -d --restart always --name sentry-cron -e
SENTRY_URL_PREFIX='http://54.152.37.40:8090' -e
SENTRY_SECRET_KEY='+i78z!+x*t*(e#qrpq#yojt5fr#td0ztv#sd!&uiwg!gcuy706'
-e SENTRY_REDIS_HOST='10.0.1.46' -e
SENTRY_POSTGRES_HOST='10.0.1.46' -e SENTRY_POSTGRES_PORT='5432' -
e SENTRY_DB_NAME='nextsentry' -e SENTRY_DB_USER='sentryapp' -e
SENTRY_DB_PASSWORD='Gh75Dc' -e
SENTRY_EMAIL_HOST='mx.learnnext.com' -e SENTRY_EMAIL_PORT='25' -e
SENTRY_SERVER_EMAIL='system@nexteducation.in' sentry run cron
```

#———— SENTRY WORKER —————

```
docker run -d --restart always --name sentry-worker-1 -e
SENTRY_URL_PREFIX='http://54.152.37.40:8090' -e
SENTRY_SECRET_KEY='+i78z!+x*t*(e#qrpq#yojt5fr#td0ztv#sd!&uiwg!gcuy706'
-e SENTRY_REDIS_HOST='10.0.1.46' -e
SENTRY_POSTGRES_HOST='10.0.1.46' -e SENTRY_POSTGRES_PORT='5432' -
e SENTRY_DB_NAME='nextsentry' -e SENTRY_DB_USER='sentryapp' -e
SENTRY_DB_PASSWORD='Gh75Dc' -e
SENTRY_EMAIL_HOST='mx.learnnext.com' -e SENTRY_EMAIL_PORT='25' -e
SENTRY_SERVER_EMAIL='system@nexteducation.in' sentry run worker
```

Things to improve :

1. Add a docker file with environment variables in it.
2. Build a docker image and run it with port mapping 8090:9000
3. Handle database(PostgreSQL) and redis(cache)

—

## Leave a comment

Comment

**Post Comment**

